



School of Computing

CS3223 Database Systems Implementation

AY21/22 Semester 2

Project Report

Team 37

Prof Tan Kian-Lee

Team Members	Student No.	Email
Kelvin Wong	A0201706U	e0415515@u.nus.edu
Moon Geonsik	A0210908L	e0484312@u.nus.edu
Wincent Tjoi	A0201480W	e0412905@u.nus.edu

Github link: <https://github.com/Wincenttjoi/CS3223-DBMS-Project>

Bonuses

Bonus Feature 1: Range Search for Mergejoin, IndexJoin and IndexSelect

Note:

Implement range search for MergeJoin, IndexJoin and IndexSelect

- MergeJoin: All non-equi operators allowed except "!="
- IndexJoin and IndexSelect:
 - If indextype is hash, only use index select scan when operator is "="
 - If indextype is btree and operator is "!=" or "<>", use normal product scan

File Changes	Changes
Planner Related Changes	
SimpleDBEngine/src/simpliedb/opt/ TablePlanner.java	<ul style="list-style-type: none">- updated makeIndexJoin() to allow non-equality search if the index supports range search. If it does, also ensure that it is not "!=" (or "<>").

	<ul style="list-style-type: none"> - updated makeMergeJoin() to allow non-equality search and rearrange parameters in MergeJoinPlan based on lhs vs rhs. Also, flip operator if ">" and ">=" to "<" and "<=" respectively (explained in MergeJoinScan) - updated makeIndexSelect() to allow non-equality search if the index supports range search. If it does, also ensure that it is not "!=" (or "<>").
SimpleDBEngine/src/simpliedb/opt/ MergeJoinScan.java	<ul style="list-style-type: none"> - Renamed the default next() method to nextEquals() - Created a method named nextLesser() that handles inequality terms e.g. lhs < rhs or lhs <= rhs. It makes use of the fact that if both scans s1 and s2 are sorted, then if lhs1 < rhs1 is false, then lhs2 < rhs1 is also false. Thus, the position of s2 is saved when the inequality term holds true for a lhs's first time. Subsequent lhs comparisons will start from the saved position of s2. - Created a method named next() that returns nextEquals or nextLesser depending on the operator
Index Related Changes	
SimpleDBEngine/src/simpliedb/index/ IndexInfo.java	<ul style="list-style-type: none"> - added supportRangeSearch(String opr) method to return true if it is supported by the idxtype
SimpleDBEngine/src/simpliedb/index/ Index.java SimpleDBEngine/src/simpliedb/index/btree/ BTPage.java SimpleDBEngine/src/simpliedb/index/btree/ BTreeDir.java	<ul style="list-style-type: none"> - overload beforeFirst(), findSlotBefore(), findChildBlock(), search() method to allow search with specific operator

SimpleDBEngine/src/simpliedb/index/hash/ H ashIndex.java	
SimpleDBEngine/src/simpliedb/index/btree/ BTreeIndex.java	<ul style="list-style-type: none"> - fix searchCost() method of BTreeIndex to return 1 when numblocks == 0. The original code would have resulted in arithmetic error with an empty table.
SimpleDBEngine/src/simpliedb/index/btree/ BTreeLeaf.java	<ul style="list-style-type: none"> - modify next() and tryOverflow() method of BTreeLeaf to check when key matches via OprComparator instead of “=”
SimpleDBEngine/src/simpliedb/index /planner/ IndexSelectPlan.java SimpleDBEngine/src/simpliedb/index/query/ I ndexSelectScan.java	<ul style="list-style-type: none"> - added operator parameter in constructors - modified beforeFirst() method of IndexSelectScan to utilize the operator in the index's beforeFirst().
Query Related Changes	
SimpleDBEngine/src/simpliedb/materialize/ O prComparator.java	<ul style="list-style-type: none"> - created a static compare method to compare 2 Constant with a given operator.
SimpleDBEngine/src/simpliedb/query/ Predicate.java	<ul style="list-style-type: none"> - add getOperatorFromConstantComparison(String fldname) and getOperatorFromFieldComparison(String fldname)
SimpleDBEngine/src/simpliedb/query/ Term.java	<ul style="list-style-type: none"> - added getOperator() method

Bonus Feature 2: Implementing a better query optimizer

Note:

Bonus 2a

The planner now considers all 4 possible implementations of join and chooses the one with lowest blocks accessed instead of simply choosing in the order of index join, hash join, merge join as suggested in Heuristic 7 of the textbook (default implementation of SimpleDB).

Bonus 2b

By default, SimpleDB's Heuristic Query Planner after creating index select plan and other join algorithms, it will create a Select Plan again on top of the same terms in the predicate.

For example, with majorid as a btree index, the query "select majorid, sid from student where majorid > 10 and sid = 4" will produce a plan of Index Select Plan on majorid > 10, and then Select Plan on majorid > 10 and sid = 4.

After our implementation to fix this, the plan used are "Index Select Plan on majorid > 10, Select Plan on sid = 4", where repetitions are avoided.

See: [Feature 2: Before and after changes](#)

Bonus 2c

After implementing bonus 2b, a query with multiple indexes in a single query would return the wrong result.

For example, if we have enroll (studentid) and student (sid) hash indexes, the query of “select studentid, sid from student, enroll where studentid = 1 and sid = 1” would return the wrong result, where sid attribute does not get filtered at all.

We further implement this to ensure that SimpleDB is working well as before.

File Changes	Changes
Planner Related Changes	
SimpleDBEngine/src/simpliedb/opt/ TablePlanner.java	<ul style="list-style-type: none">- added class variables selectTermToRemove and joinTermsSelectedToRemove, which store the terms to be excluded when returning plans for index select and corresponding join algorithms.- makeSelectPlan will return a plan with all applicable select terms, without any select term used by makeIndexSelect. This is done while maintaining the selected term in the TablePlanner's predicate so that the TablePlanner can be reused.

	<ul style="list-style-type: none"> - makeJoinPlan will return a plan with all applicable join and select terms, without any joinTerm used by the join algorithm while maintaining the joinTerm in the TablePlanner's predicate. This is done while maintaining the joinTerm in the TablePlanner's predicate so that the TablePlanner can be reused. - modified makeJoinPlan routine to create all join plans, then sorting based on blocks accessed and choosing the plan with least estimated blocks accessed.
Query Related Changes	
SimpleDBEngine/src/simpliedb/query/ Term.java	<ul style="list-style-type: none"> - added matches(Term) method to check if this term matches another term, regardless of direction. E.g. $(a > 3)$ matches $(3 < a)$ and $(c \geq d)$ matches $(d \geq c)$
SimpleDBEngine/src/simpliedb/query/ Predicate.java	<ul style="list-style-type: none"> - added new method removeTerm(Term) to remove all terms that match the parameter

Bonus Feature 3: Float data type which supports group by AVG

File Changes	Changes
Planner Related Changes	
SimpleDBEngine/src/simpliedb/materialize/ GroupByPlan .java	<ul style="list-style-type: none">- refactored GroupByPlan constructor to add a float field to the schema for average aggregation function
SimpleDBEngine/src/simpliedb/record/ Schema .java	<ul style="list-style-type: none">- created addFloatField() to support float data type field in schema
SimpleDBEngine/src/simpliedb/jdbc/.../ ResultSet .java (etc.) SimpleDBEngine/src/simpliedb/.../ Plan .java (etc.) SimpleDBEngine/src/simpliedb/.../ Scan .java (etc.)	<ul style="list-style-type: none">- refactored to support float data type e.g. implement getFloat(), setFloat(), etc.- applies to any class that implements getInt()/setString() and getString()/setString()
SimpleDBEngine/src/simpliedb/query/ Constant .java	<ul style="list-style-type: none">- refactored to hold float value and support comparison and hashing of float values
SimpleDBEngine/src/simpliedb/materialize/ AvgFn .java	<ul style="list-style-type: none">- refactored to hold float value and return Constant initialised with float value

Bonus Feature 4: Joining an empty table with no tuples to another table

Note:

By default, SimpleDB will throw an error if a table with no tuples is joined with another table.

We have handled this gracefully that we will return an empty result to the user, since conceptually such a query can exist and a user wants to know that there is an empty result.

See: [Feature 4: Before and after changes](#)

File Changes	Changes
Planner Related Changes	
SimpleDBEngine/src/simpliedb/materialize/ MaterializedPlan .java	<ul style="list-style-type: none">- in open(), we keep track of boolean isTupleExist, only when src.next() is true, isTupleExist is toggled to true. At the end of open(), isTupleExist will be checked, and if true, setTupleExist to be true in dest temporary table.
SimpleDBEngine/src/simpliedb/multibuffer/ MultibufferProductScan .java	<ul style="list-style-type: none">- next() will have an additional guard clause to check if lhsscan have a tuple. If no tuple, return false.

SimpleDBEngine/src/simpliedb/record/
TableScan.java

- new class field of **isTupleExist** (boolean) for tracking purposes
- upon construction of object, set **isTupleExist** to be false
- when **next()** is called and returns true, set **isTupleExist** to true
- added setter and getter for **isTupleExist**

Labs

Lab 2: Support for hash index and B+-tree index

File Changes	Changes
Index and Metadata	
SimpleDBEngine/src/simplydb/metadata/ IndexInfo.java	<ul style="list-style-type: none">- stored idxtype as private field- new public setter method setIdxType to set private index- open() and blocksAccessed() to call hashindex / btreeindex depending on idxtype field of IndexInfo
SimpleDBEngine/src/simplydb/metadata/ IndexMgr.java	<ul style="list-style-type: none">- schema to add another string field of "idxtype" for tblmgr- createIndex takes in new parameter idxtype, which the tableScan ts will setString on "idxtype"- getIndexInfo have getString("idxtype") which will be used to setIdxType for IndexInfo
SimpleDBEngine/src/simplydb/metadata/ MetadataMgr.java	<ul style="list-style-type: none">- createIndex have new parameter String idxtype, and idxmgr.createIndex will pass the idxtype
SimpleDBEngine/src/simplydb/metadata/ MetadataMgrTest.java	<ul style="list-style-type: none">- mdm.createIndex to support hardcoded "hash"/"btree" as new parameters
SimpleDBEngine/src/simplydb/parse/ CreateIndexData.java	<ul style="list-style-type: none">- supported new private field idxtype and updated constructor- getter method for indexType()

Parser Related Changes

SimpleDBEngine/src/simpliedb/parse/
Lexer.java

- new private field of Collection<String> **idxType**
- **initIdxType()** which add keywords of “hash” and “btree”, initialized during **Lexer** object creation
- added “using” inside **initKeywords()**
- implemented **matchIdxType** and **eatIdxType**, similar to other match and eat methods

SimpleDBEngine/src/simpliedb/parse/
Parser.java

- **createIndex()** will **eatKeyword**(“using”)
- **createIndexData** takes in new parameter **idxtype**, which is retrieved by **lex.eatIdxType()**

Planner Related Changes

SimpleDBEngine/src/simpliedb/index/
planner/**IndexUpdatePlanner.java**

- **executeCreateIndex** to call **mdm.createIndex** with a new parameter **indexType**

SimpleDBEngine/src/simpliedb/plan/
BasicUpdatePlanner.java

- **mdm.createIndex** to take in new parameter **data.indexType()**

SimpleDBEngine/src/simpliedb/server/
SimpleDB.java

- used **HeuristicQueryPlanner** and **IndexUpdatePlanner** instead of basic planners

Planner Related Changes

CreateStudentDB.java	- added create index statements on student / enroll with the keyword “using hash/btree”
-----------------------------	---

See also: [Lab 2: Sequence Diagram for IdxMgr and TableScan](#)

Lab 3: Support for order by clause and sorting

File Changes	Changes
Parser Related Changes	
SimpleDBEngine/src/simplydb/parse/ Lexer.java	<ul style="list-style-type: none">- added keywords “order” and “by” to the keywords list (in routine initKeywords())- created a list of sort type (“asc,” “desc”) (in routine initSortType())- added a new routine eatSortType() to take in the sort type and return a boolean value (“asc”=>true; “desc”=>false;)- added a new routine matchSortType() to check if the current token is a legal sort type (“asc,” “desc”)
SimpleDBEngine/src/simplydb/parse/ Parser.java	<ul style="list-style-type: none">- modified routine query() to detect the “order by” keyword and call sortList(). Throw a BadSyntaxException, if invalid “order by” syntax is used- added a new routine sortList() (called when the “order” keyword is detected) to parse the “order by” clause and to initialize and return a LinkedHashMap<String, Boolean> containing <key, value> pairs of the field to be sorted and whether it should be sorted in an ascending or descending order (order is set to ascending by default)
Planner Related Changes	

SimpleDBEngine/src/simpliedb/opt/ HeuristicQueryPlanner.java	<ul style="list-style-type: none"> - modified the routine createPlan() to initialize and return a SortPlan object, using the ProjectPlan object, if the query contains an “order by” clause. Otherwise, return a ProjectPlan object
Sorting Related Changes	
SimpleDBEngine/src/simpliedb/materialize/ SortPlan.java	<ul style="list-style-type: none"> - modified the constructor of SortPlan object that iterates over sortfields to assign a true Boolean value to each field as a key to create a Map<String, Boolean> object and initialize a RecordComparator object with the Map object - created a new constructor for SortPlan object that takes in a Map<String, Boolean> object, instead of a List<String> object, and initialize a RecordComparator object with the Map object
SimpleDBEngine/src/simpliedb/materialize/ RecordComparator.java	<ul style="list-style-type: none"> - modified the constructor of RecordComparator object to take in a Map<String, Boolean> object and initialize the Map object as sortMap. Also, extract the keyset of the Map object as a List<String> and initialize the List object as fields - modified the routine compare() to check if the current field that the routine is iterating over should be sorted in an “ascending” or “descending” order using the sortMap and return the integer result. (if the order is descending, change the sign of the integer result and return)

Lab 4: Supporting for nested-loops join, sort-merge join and index-based join

File Changes	Changes
Parser Related Changes	
SimpleDBEngine/src/simpliedb/parse/ Parser.java	<ul style="list-style-type: none">- added ability to parse keywords ‘join’ and ‘on’ in tableList() and extract the table name and predicate
Planner Related Changes	
SimpleDBEngine/src/simpliedb/opt/ TablePlanner.java	<ul style="list-style-type: none">- added makeMergeJoin routine to create MergeJoinPlan- added makeNestedJoin routine to create NestedJoinPlan- modified makeJoinPlan routine to allow either indexJoinPlan, mergeJoinPlan or nestedJoinPlan to be created based on heuristic 7 of the textbook
SimpleDBEngine/src/simpliedb/materialize/ NestedJoinScan.java	<ul style="list-style-type: none">- Created nestedJoinPlan and NestedJoinScan based on ProductPlan, and checking the predicate for each pair of records.
SimpleDBEngine/src/simpliedb/materialize/ NestedJoinPlan.java	
Sorting Related Changes	

SimpleDBEngine/src/simpliedb/materialize/ SortPlan.java	<ul style="list-style-type: none"> - modified condition in open() to end generation of sorted runs when size ≤ 1 instead of 2. - added blocksAccessedWithSortCost() method to estimate number of block accesses for sorting and materialization with the formula taught in syllabus
SimpleDBEngine/src/simpliedb/materialize/ MergeJoinPlan.java	<ul style="list-style-type: none"> - updated blocksAccessed() method to include the sorting cost of both plans using SortPlan.blocksAccessedWithSortCost(), assuming best case scenario
Query Related Changes	
SimpleDBEngine/src/simpliedb/query/ Constant.java	<ul style="list-style-type: none"> - equals() returns false if obj being compared is null
SimpleDBEngine/src/simpliedb/query/ Predicate.java	<ul style="list-style-type: none"> - added get method for list of terms
SimpleDBEngine/src/simpliedb/query/ Term.java	<ul style="list-style-type: none"> - added get method for LHS and RHS expressions

Lab 5: Support for partition-based (or hash) join, and aggregates (SUM, COUNT, AGV, MIN, MAX) with/without group by clause (using sort-based implementation of group by operator)

File Changes	Changes
Parser Related Changes	
SimpleDBEngine/src/simpliedb/parse/ Parser.java	<ul style="list-style-type: none">- created aggField() to extract information from the aggregate field- refactored selectList() to parse the select fields and properly extract information about the aggregation function from the aggregate fields- created parseAggFn() to initialize and return the correct AggregationFunction- created groupList() to extract information from the group by fields- refactored query() to parse aggregate fields and group by fields
SimpleDBEngine/src/simpliedb/parse/ Lexer.java	<ul style="list-style-type: none">- added keywords “group” and “hashjoin” to the keywords list (in routine initKeywords())- created a list of aggregation type (“sum”, “count”, “avg”, “min,” “max”) (in routine initAggType())- added a new routine eatAggType() to take in the aggregation type and return a string value- added a new routine matchAggType() to check if the current token is a legal aggregation type

SimpleDBEngine/src/simpliedb/parse/ QueryData.java	<ul style="list-style-type: none"> - refactored QueryData to initialize a list of aggregation functions and a list of group by fields - refactored toString() to print out information about aggregate fields and group by fields
Group By Related Changes	
SimpleDBEngine/src/simpliedb/opt/ HeuristicQueryPlanner.java	<ul style="list-style-type: none"> - refactored createPlan() to aggregate results on group fields, if either aggregate fields or group by fields are included in the query
SimpleDBEngine/src/simpliedb/materialize/ AvgFn.java	<ul style="list-style-type: none"> - created AvgFn class to compute the average of values in the same group - processFirst() initializes the sum as the value of the first record and initializes the num as 1 - processNext() accumulates the sum by the value of the current record and increments the num - value() returns a Constant object initialized with sum divided by num
SimpleDBEngine/src/simpliedb/materialize/ MinFn.java	<ul style="list-style-type: none"> - created MinFn class to compute the minimum of values in the same group - processFirst() initializes the minimum as the value of the first record - processNext() updates the minimum to be the value of the current record, if the value is less than the current minimum - value() returns a Constant object initialized with the current minimum
SimpleDBEngine/src/simpliedb/materialize/ SumFn.java	<ul style="list-style-type: none"> - created SumFn class to compute the sum of values in the same group - processFirst() initializes the sum as the value of the first record - processNext() accumulates the sum by the value of the current record - value() returns a Constant object initialized with the current sum

Hash Join Related Changes

SimpleDBEngine/src/simpliedb/opt/ TablePlanner.java	<ul style="list-style-type: none">- refactored makeJoinPlan() to initialize a HashJoinPlan and compare the number of accessed blocks with other joins- created makeHashJoin() to execute hash join, if the operator of the predicate is equality and enough buffer size is available ($B > \sqrt{M}$). Otherwise, execute product join
SimpleDBEngine/src/simpliedb/materialize/ HashComparator.java	<ul style="list-style-type: none">- created HashComparator class to check if the two scans are hashed to the same bucket, based on the values of the specified fields
SimpleDBEngine/src/simpliedb/materialize/ HashJoinPlan.java	<ul style="list-style-type: none">- created HashJoinPlan class to initialize a plan that joins two tables using hashjoin method- open() initializes a hashjoin scan and hashes the record of each table into k number of partitions (temporary tables) using splitIntoPartitions()- splitIntoPartitions() initializes k number of partitions (temporary tables) and scans the given table to hash each record into the corresponding partitions
SimpleDBEngine/src/simpliedb/materialize/ HashJoinScan.java	<ul style="list-style-type: none">- created HashJoinScan class to initialize a scan that probes two matching partitions and finds the matching records using a different hash function- next() traverses through all sets of matching partitions and returns true when a set of matching records that are hashed to the same hash table are found

Lab 6: Support for DISTINCT and displaying the query plan

File Changes	Changes
Parser Related Changes	
SimpleDBEngine/src/simpliedb/parse/ Lexer.java	<ul style="list-style-type: none">- added keywords “distinct” to the keywords list (in routine initKeywords())
SimpleDBEngine/src/simpliedb/parse/ Parser.java	<ul style="list-style-type: none">- added ability to parse keywords ‘distinct’ in query() after eating keyword ‘distinct’
Planner Related Changes	
SimpleDBEngine/src/simpliedb/opt/ HeuristicQueryPlanner.java	<ul style="list-style-type: none">- added another condition to check if QueryData isDistinct field is true, then calls sort plan and distinctplan.
SimpleDBEngine/src/simpliedb/materialize/ DistinctPlan.java	<ul style="list-style-type: none">- DistinctPlan implements Plan, which calls DistinctScan(Scan, fields) when open() is called.
SimpleDBEngine/src/simpliedb/materialize/ DistinctScan.java	<ul style="list-style-type: none">- contains private fields of Scan s, List<String> fields, and List<Constant> of prev and curr (to keep track of current and previous record values)- calls internal method nextDistinct() when next() is called

	<ul style="list-style-type: none"> - nextDistinct() returns a boolean, and calls s.next() recursively if the next record is a duplicate, returns false when no more record. Each field of the projected query has their Constant value compared, to know whether the next record is a duplicate.
Query Related Changes	
SimpleDBEngine/src/simpliedb/parse/ QueryData.java	<ul style="list-style-type: none"> - holds a new boolean field isDistinct - isDistinct field supported during QueryData object initialization - getter method for isDistinct - toString method to include isDistinct, if any
Displaying Query Related Changes	
SimpleDBEngine/src/simpliedb/index/ planner/ IndexJoinPlan.java SimpleDBEngine/src/simpliedb/index/ planner/ IndexSelectPlan.java SimpleDBEngine/src/simpliedb/materialize/ D istinctPlan.java SimpleDBEngine/src/simpliedb/materialize/ G roupByPlan.java SimpleDBEngine/src/simpliedb/materialize/ MaterializePlan.java	<ul style="list-style-type: none"> - all classes that implement Plan will have to implement printPlan() method, which tells users that the plan is being used. Specific details could also be shown when necessary. - every open() method will call printPlan() method, after opening the scan further in.

SimpleDBEngine/src/simpliedb/materialize/ MergeJoinPlan.java	
SimpleDBEngine/src/simpliedb/materialize/ S ortPlan.java	
SimpleDBEngine/src/simpliedb/multibuffer/ M ultibufferProductPlan.java	
SimpleDBEngine/src/simpliedb/plan/ OptimizedProductPlan.java	
SimpleDBEngine/src/simpliedb/plan/ ProductPlan.java	
SimpleDBEngine/src/simpliedb/plan/ ProjectPlan.java	
SimpleDBEngine/src/simpliedb/plan/ SelectPlan.java	
SimpleDBEngine/src/simpliedb/plan/ TablePlan.java	
SimpleDBEngine/src/simpliedb/plan/ Plan.java	

- interface **Plan** to have **printPlan()** method, which displays to the user that the plan is used.

See also: [Lab 6: Sequence Diagram for Distinct](#), [Lab 6: Classes implementing Plan](#)

Lab 7: Integrate all features into SimpleDB+

Note:

1. We have been continuously integrating each lab, so there is not much effort needed for integration here.
2. Add range search for nestedJoin.
3. Heuristic Planner will be used by default, Basic Planner no longer used.

File Changes	Changes
Parser Related Changes	
SimpleDBEngine/src/simpliedb/parse/ Lexer.java SimpleDBEngine/src/simpliedb/parse/ Parser.java	<ul style="list-style-type: none">- added "indexjoin", "mergejoin", "nestedjoin" and "hashjoin" as keywords- added information on selected join algorithm to data if the above keywords are found before 'select' keyword
Planner Related Changes	

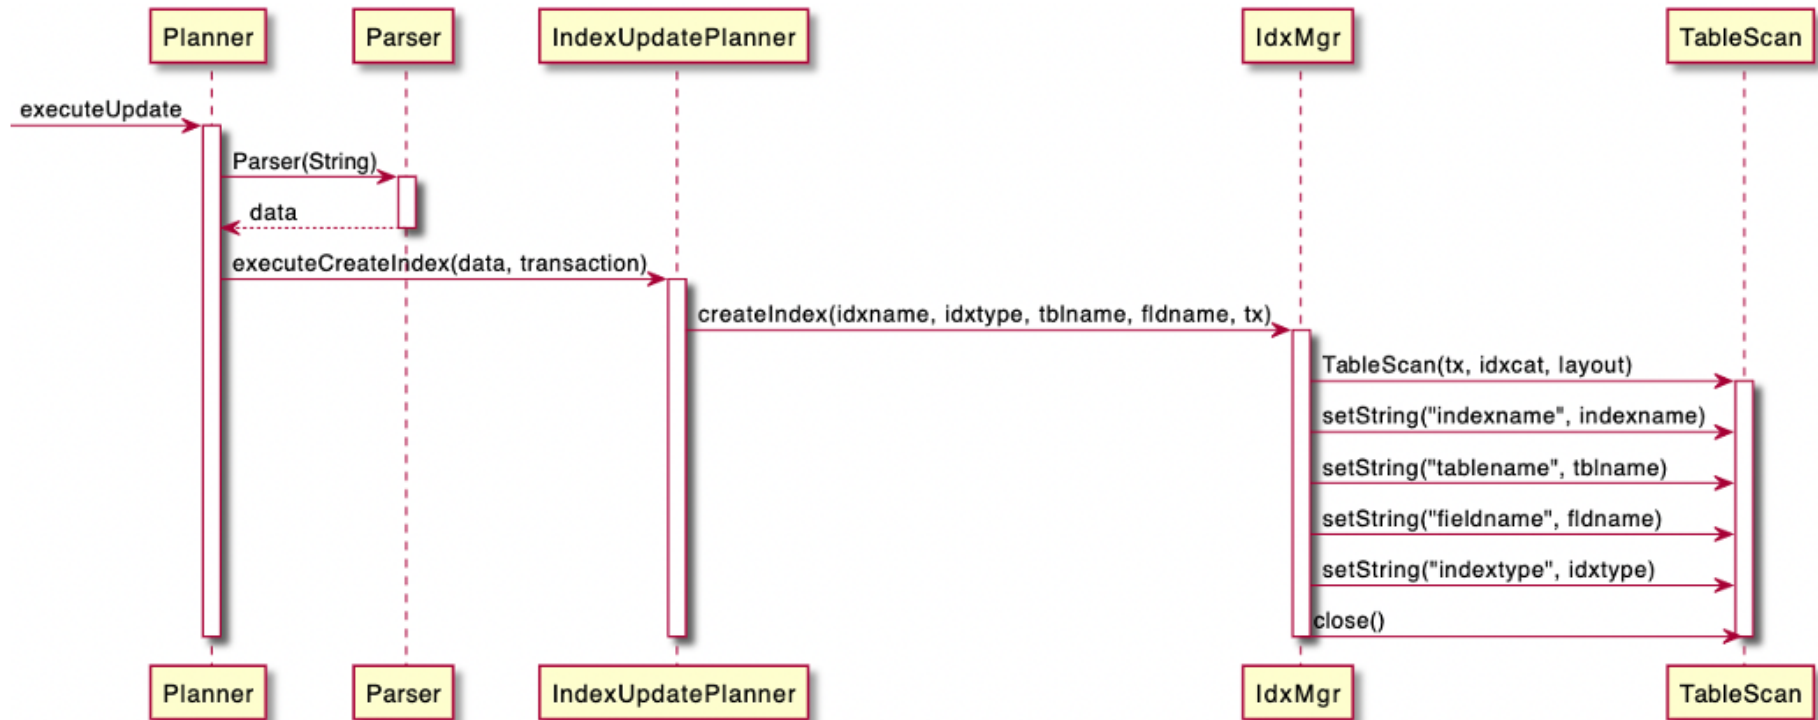
SimpleDBEngine/src/simpliedb/opt/
TablePlanner.java

SimpleDBEngine/src/simpliedb/opt/
HeuristicQueryPlanner.java

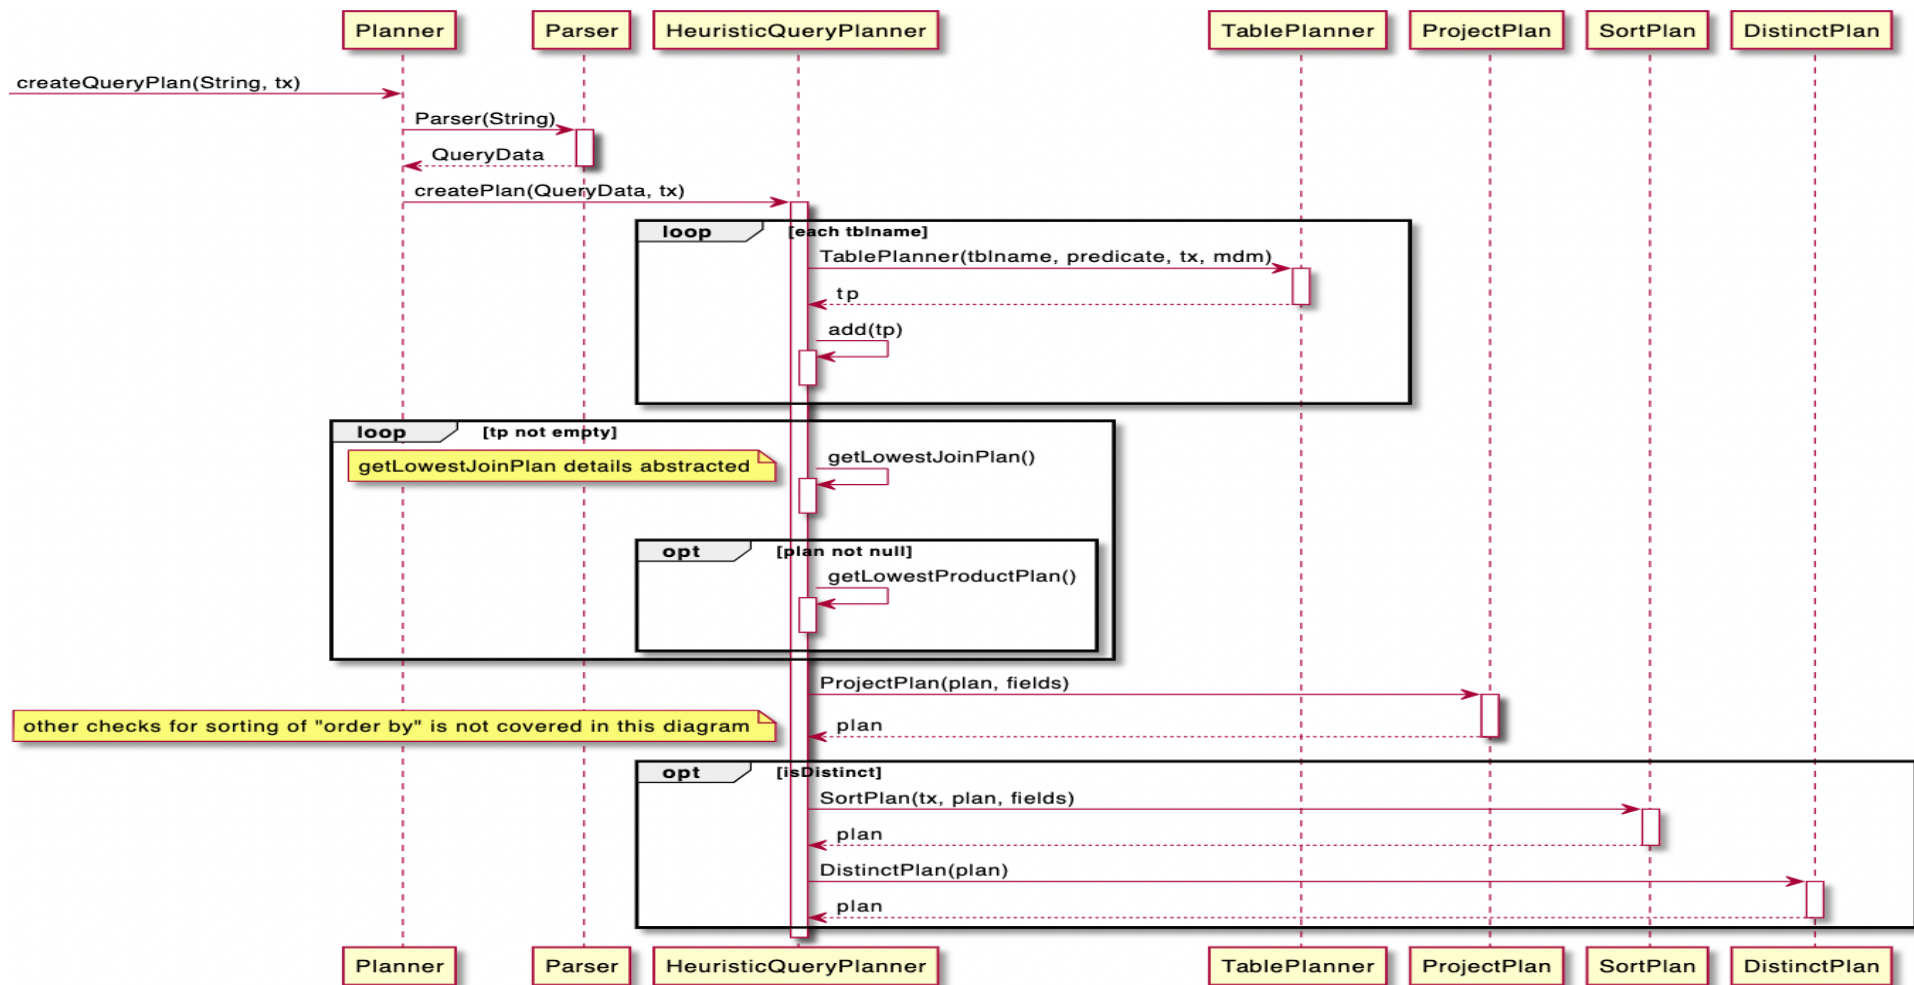
- updated to allow a specific join algorithm to be chosen for showcase in future

Appendix

Lab 2: Sequence Diagram for IdxMgr and TableScan

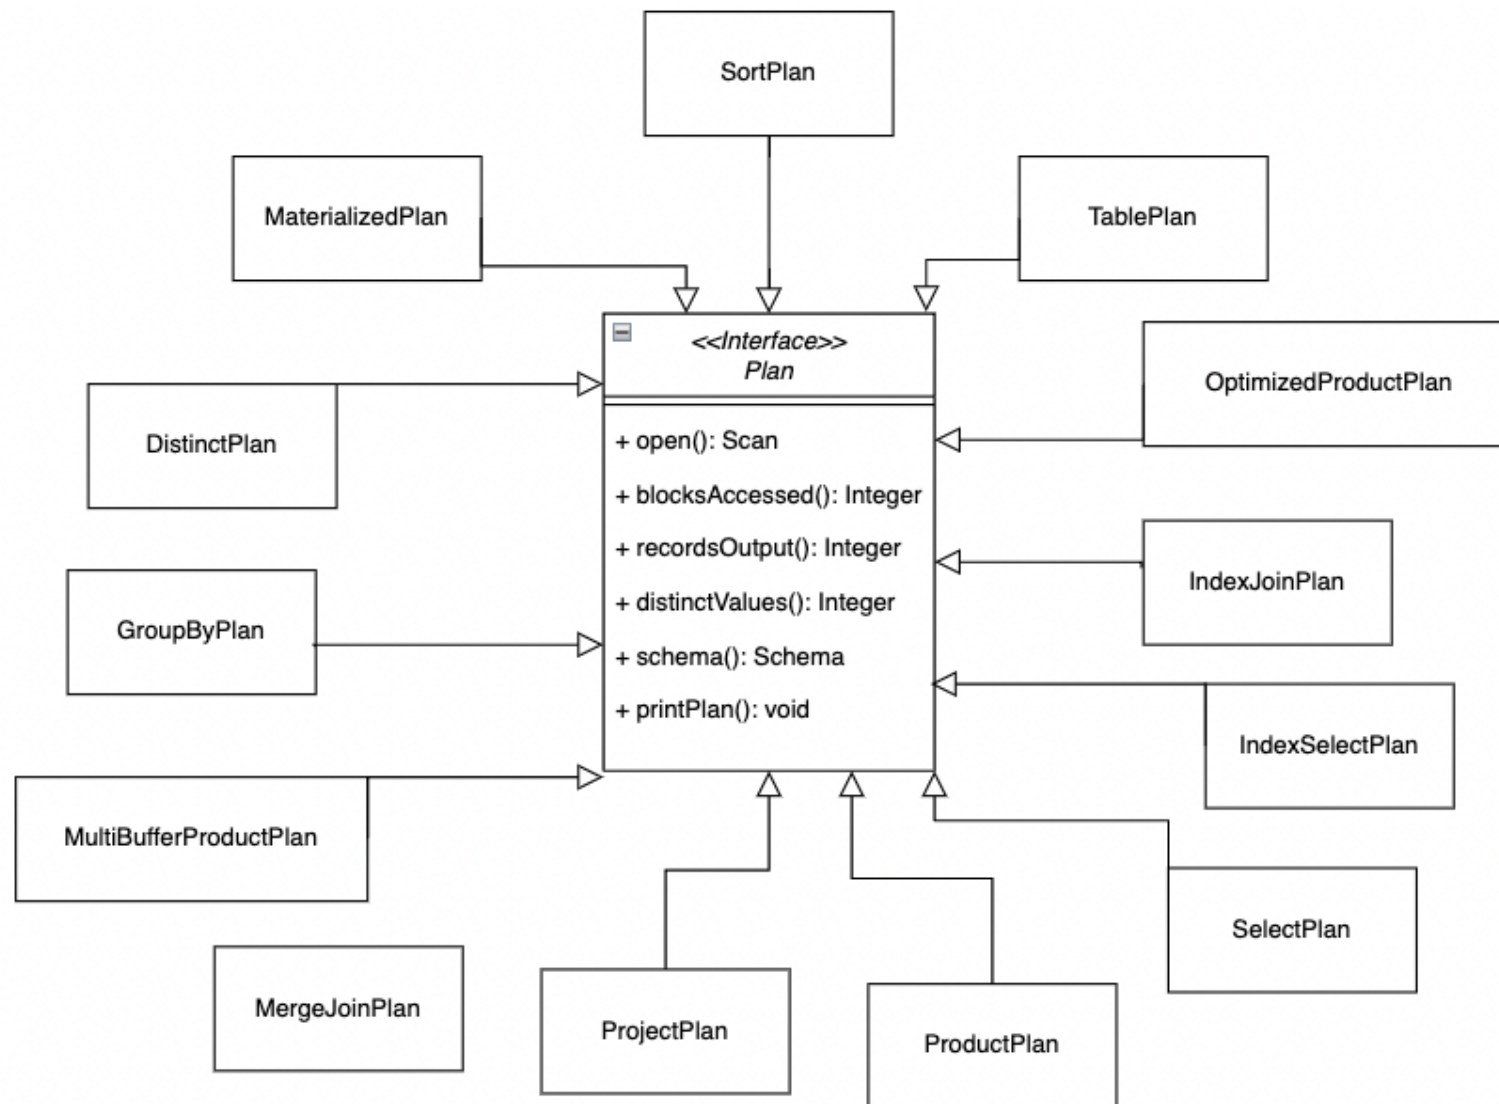


Lab 6: Sequence Diagram for Distinct



Note: HashJoinPlan, NestedJoinPlan and other plans are excluded for brevity.

Lab 6: Classes implementing Plan



Feature 2: Before and after changes

```
110
111     Test.doTest(stmt, "select sname, sid, majorid from student where majorid >= 20 and sid > 4 and sname != 'bob'");
112 // TEST9
113 //      sname      sid majorid
114 //      -----
115 //          kim        6      20
116 //          art        7      30
117 //          pat        8      20
118
```

Console × Problems Debug Shell Search

<terminated> Lab1Test (1) [Java Application] /Users/wincenntjoi/Library/Java/JavaVirtualMachines/openjdk-17.0.1/Contents/Home/bin/java (15 Mar 2022, 10:56:22 pm – 10:56:23 pm)

TEST9
Table Plan used on table student
Index Select Plan used for table student of field majorid
Select Plan used on predicate majorid>=20 and sid>4 and sname!=bob
Project Plan used on fields of sname, sid, majorid

sname	sid	majorid
pat	8	20
kim	6	20
art	7	30

The total time taken for query is: 747375ns
transaction 10 committed

IndexSelect: Before


```

TEST9: select sname, sid, majorid from student where majorid >= 20 and sid > 4 and sname != 'bob'
Table Plan used on table student
Index Select Plan used for table student of field majorid>=20, index accessed: majorid_idx, index type: btree
Select Plan used on predicate sid>4 and sname!=bob
Project Plan used on fields of sname, sid, majorid
The total time taken for query is: 400583ns
  sname    sid majorid
-----
    pat      8     20
    kim      6     20
    art      7     30
transaction 10 committed
=====

```

IndexSelect: After

```

TEST4: indexjoin select sname, sid, studentid, grade from student, enroll where studentid = sid
>>>Indexjoin blocks accessed = 13
Table Plan used on table enroll
Table Plan used on table student
Index join Plan used on sid=studentid
Select Plan used on predicate studentid=sid
Project Plan used on fields of sname, sid, studentid, grade
The total time taken for query is: 3ms

```

sname	sid	studentid	grade
joe	1	1	A
joe	1	1	C
amy	2	2	B+
sue	4	4	B
sue	4	4	A
kim	6	6	A

transaction 11 committed

IndexJoin: Before

```

TEST4: indexjoin select sname, sid, studentid, grade from student, enroll where studentid = sid
>>>Indexjoin blocks accessed = 13
Table Plan used on table enroll
Table Plan used on table student
Index join Plan used on sid=studentid
Project Plan used on fields of sname, sid, studentid, grade
The total time taken for query is: 2ms

```

sname	sid	studentid	grade
joe	1	1	A
joe	1	1	C
amy	2	2	B+
sue	4	4	B
sue	4	4	A
kim	6	6	A

IndexJoin: After

Feature 4: Before and after changes

```
TEST17: select sname from student, dept where sname = 'jonah'
Table Plan used on table student
Select Plan used on predicate sname=jonah
Materialized Plan created with fields [sid, sname, majorid, gradyear]
Table Plan used on table dept
Multibuffer Product Plan used between simpledb.materialize.MaterializePlan@6d311334 and simpledb.plan.TablePlan@682a0b20
Project Plan used on fields of sname
The total time taken for query is: 2ms
      sname
-----
transaction 18 rolled back
SQL Exception: java.lang.IllegalArgumentException: newPosition < 0: (-22 < 0)
```

Joining with an empty table : Before

```
TEST17: select sname from student, dept where sname = 'jonah'
Table Plan used on table student
Select Plan used on predicate sname=jonah
Materialized Plan created with fields [sid, sname, majorid, gradyear]
Table Plan used on table dept
Multibuffer Product Plan used between simpledb.materialize.MaterializePlan@682a0b20 and simpledb.plan.TablePlan@3d075dc0
Project Plan used on fields of sname
The total time taken for query is: 2ms
      sname
-----
transaction 18 committed
=====
```

Joining with an empty table : After